

```
In [1]: import Pkg
Pkg.activate(@_DIR_)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
using Test
import Convex as cvx
import ECOS
using Random

Activating environment at `~/ocrl_ws/16745-ocrl/HW2_S23/Project.toml`
Updating registry at `~/julia/registries/General`
Installed Scratch _____ v1.2.0
Installed Adapt _____ v3.6.1
Installed GenericLinearAlgebra _ v0.3.7
Installed Parsers _____ v2.5.8
Installed Compat _____ v4.6.1
Installed JLD2 _____ v0.4.31
Installed ForwardDiff _____ v0.10.35
Installed Plots _____ v1.38.7
Installed MathOptInterface _____ v1.13.1
Installed IrrationalConstants _ v0.2.2
Updating `~/ocrl_ws/16745-ocrl/HW2_S23/Project.toml`
[f65535da] + Convex v0.15.3
[e2685f51] + ECOS v1.1.0
[f6369f11] + ForwardDiff v0.10.35
[b99e6be6] + Hypatia v0.7.0
[033835bb] + JLD2 v0.4.31
[283c5d60] + MeshCat v0.14.2
[015f5b4d] + Plots v1.38.7
```

Julia Warnings:

- 1. For a function `foo(x::Vector)` with 1 input argument, it is not neccessary to do `df_dx = FD.jacobian(_x -> foo(_x), x)` . Instead you can just do `df_dx = FD.jacobian(foo, x)` . If you do the first one, it can dramatically slow down your compilation time.
- 2. Do not define functions inside of other functions like this:

```
function foo(x)
    # main function foo

    function body(x)
        # function inside function (DON'T DO THIS)
        return 2*x
    end

    return body(x)
end
```

This will also slow down your compilation time dramatically.

Q1: Finite-Horizon LQR (50 pts)

For this problem we are going to consider a "double integrator" for our dynamics model. This system has a state  $x \in \mathbb{R}^4$ , and control  $u \in \mathbb{R}^2$ , where the state describes the 2D position  $p$  and velocity  $v$  of an object, and the control is the acceleration  $a$  of this object. The state and control are the following:

$$x = [p_1, p_2, v_1, v_2]$$
$$u = [a_1, a_2]$$

And the continuous time dynamics for this system are the following:

$$\dot{x} = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} x + \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 1 & 0 \\ 0 & 1 \end{bmatrix} u$$

Part A: Discretize the model (5 pts)

Use the matrix exponential ( `exp` in Julia) to discretize the continuous time model. See the [first recitation \(https://youtu.be/EjAiRam95U4\)](https://youtu.be/EjAiRam95U4) if you're unsure of what to do.

```
In [2]: # double integrator dynamics
function double_integrator_AB(dt)::Tuple{Matrix,Matrix}
    Ac = [0 0 1 0;
          0 0 0 1;
          0 0 0 0;
          0 0 0 0.]
    Bc = [0 0;
          0 0;
          1 0;
          0 1]
    nx, nu = size(Bc)

    # TODO: discretize this linear system using the Matrix Exponential

    A = zeros(nx,nx) # TODO
    B = zeros(nx,nu) # TODO

    expmat = exp([Ac Bc; zeros(nu, nx + nu)] .* dt)

    A = expmat[1:nx,1:nx]
    B = expmat[1:nx,nx+1:nx+nu]

    @assert size(A) == (nx,nx)
    @assert size(B) == (nx,nu)

    return A, B
end

Out[2]: double_integrator_AB (generic function with 1 method)

In [3]: @testset "discrete time dynamics" begin
    dt = 0.1
    A,B = double_integrator_AB(dt)

    x = [1,2,3,4.]
    u = [-1,-3.]

    @test isapprox((A*x + B*u),[1.295, 2.385, 2.9, 3.7];atol = 1e-10)
end

Test Summary: | Pass Total
discrete time dynamics | 1 1

Out[3]: Test.DefaultTestSet("discrete time dynamics", Any[], 1, false, false)
```

Part B: Finite Horizon LQR via Convex Optimization (15 pts)

We are now going to solve the finite horizon LQR problem with convex optimization. As we went over in class, this problem requires  $Q \in \mathcal{S}_+$  (Q is symmetric positive semi-definite) and  $R \in \mathcal{S}_{++}$  (R is symmetric positive

definite). With this, the optimization problem can be stated as the following:

$$\begin{aligned} \min_{x_1:N,u_1:N-1} \quad & \sum_{i=1}^{N-1} \left[ \frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = A x_i + B u_i \quad \text{for } i = 1, 2, \dots, N - 1 \end{aligned}$$

This problem is a convex optimization problem since the cost function is a convex quadratic and the constraints are all linear equality constraints. We will setup and solve this exact problem using the `Convex.jl` modeling package. (See 2/16 Recitation video for help with this package. [Notebook is here \(https://github.com/Optimal-Control-16-745/recitations/blob/main/2\\_17\\_recitation/Convex.jl\\_tutorial.ipynb\)](https://github.com/Optimal-Control-16-745/recitations/blob/main/2_17_recitation/Convex.jl_tutorial.ipynb).) Your job in the block below is to fill out a function `Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x_ic)` , where you will form and solve the above optimization problem.

```
In [32]: # utilities for converting to and from vector of vectors <-> matrix
function mat_from_vec(X::Vector{Vector{Float64}})::Matrix
    # convert a vector of vectors to a matrix
    Xm = hcat(X...)
    return Xm
end
function vec_from_mat(Xm::Matrix)::Vector{Vector{Float64}}
    # convert a matrix into a vector of vectors
    X = [Xm[:,i] for i = 1:size(Xm,2)]
    return X
end
"""
X,U = convex_trajopt(A,B,Q,R,Qf,N,x_ic; verbose = false)

This function takes in a dynamics model x_{k+1} = A*x_k + B*u_k
and LQR cost Q,R,Qf, with a horizon size N, and initial condition
x_ic, and returns the optimal X and U's from the above optimization
problem. You should use the `vec_from_mat` function to convert the
solution matrices from cvx into vectors of vectors (vec_from_mat(X.value))
"""
function convex_trajopt(A::Matrix,      # A matrix
                       B::Matrix,      # B matrix
                       Q::Matrix,      # cost weight
                       R::Matrix,      # cost weight
                       Qf::Matrix,     # term cost weight
                       N::Int64,       # horizon size
                       x_ic::Vector;    # initial condition
                       verbose = false
                       )::Tuple{Vector{Vector{Float64}},Vector{Vector{Float64}}}

    # check sizes of everything
    nx,nu = size(B)
    @assert size(A) == (nx, nx)
    @assert size(Q) == (nx, nx)
    @assert size(R) == (nu, nu)
    @assert size(Qf) == (nx, nx)
    @assert length(x_ic) == nx

    # TODO:
    # create cvx variables where each column is a time step
    # hint: x_k = X[:,k], u_k = U[:,k]
    X = cvx.Variable(nx, N)
    U = cvx.Variable(nu, N - 1)

    # create cost
    # hint: you can't do x'*Q*x in Convex.jl, you must do cvx.quadform(x,Q)
    # hint: add all of your cost terms to `cost`
    cost = 0
    for k = 1:(N-1)

        # add stagewise cost
        xk = X[:,k]
        uk = U[:,k]
        cost += (0.5 * cvx.quadform(xk,Q))
        cost += (0.5 * cvx.quadform(uk,R))
    end

    # add terminal cost
    xn = X[:,N]
    cost += (0.5 * cvx.quadform(xn,Qf))

    # initialize cvx problem
    prob = cvx.minimize(cost)

    # TODO: initial condition constraint
    # hint: you can add constraints to our problem like this:
    # prob.constraints += (Gz == h)
    prob.constraints += (X[:,1] == x_ic)
    for k = 1:(N-1)
        # dynamics constraints
        prob.constraints += (X[:,k+1] == A*X[:,k] + B*U[:,k])
    end
    # solve problem (silent solver tells us the output)
    cvx.solve!(prob, ECOS.Optimizer; silent_solver = !verbose)

    if prob.status != cvx.MathOptInterface.OPTIMAL
        error("Convex.jl problem failed to solve for some reason")
    end

    # convert the solution matrices into vectors of vectors
    X = vec_from_mat(X.value)
    U = vec_from_mat(U.value)

    return X, U
end
```

```
Out[32]: convex_trajopt
```

Now let's solve this problem for a given initial condition, and simulate it to see how it does:

In [34]: @testset "LQR via Convex.jl" begin

```
# problem setup stuff
dt = 0.1
tf = 5.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 5*Q

# initial condition
x_ic = [5,7,2,-1.4]

# setup and solve our convex optimization problem (verbose = true for submission)
Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x_ic; verbose = false)

# TODO: simulate with the dynamics with control Ucvx, storing the
# state in Xsim

# initial condition
Xsim = [zeros(nx) for i = 1:N]
Xsim[1] = 1*x_ic



# TODO dynamics simulation
for k = 1:(N-1)
    # dynamics constraints
    Xsim[k+1] = A * Xsim[k] + B * Ucvx[k]
end

@test length(Xsim) == N
@test norm(Xsim[end])>1e-13
#-----plotting-----
Xsim_m = mat_from_vec(Xsim)

# plot state history
display(plot(t_vec,Xsim_m',label = ["x₁" "x₂" "ẋ₁" "ẋ₂"],
            title = "State History",
            xlabel = "time (s)", ylabel = "x"))

# plot trajectory in x1 x2 space
display(plot(Xsim_m[1,:],Xsim_m[2,:],
            title = "Trajectory in State Space",
            ylabel = "x₂", xlabel = "x₁", label = ""))
#-----plotting-----

# tests
@test 1e-14 < maximum(norm.(Xsim .- Xcvx,Inf)) < 1e-3
@test isapprox(Ucvx[1], [-7.8532442316767, -4.127120137234], atol = 1e-3)
@test isapprox(Xcvx[end], [-0.02285990, -0.07140241, -0.21259, -0.1540299], atol = 1e-3)
@test 1e-14 < norm(Xcvx[end] - Xsim[end]) < 1e-3
end
```



Test Summary: | Pass Total  
LQR via Convex.jl | 6 6

Out[34]: Test.DefaultTestSet("LQR via Convex.jl", Any[], 6, false, false)

Bellman's Principle of Optimality

Now we will test Bellman's Principle of optimality. This can be phrased in many different ways, but the main gist is that any section of an optimal trajectory must be optimal. Our original optimization problem was the above problem:

$$\begin{aligned} \min_{x_{1:N}, u_{1:N-1}} \quad & \sum_{i=1}^{N-1} \left[ \frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = A x_i + B u_i \quad \text{for } i = 1, 2, \dots, N-1 \end{aligned}$$

which has a solution  $x_{1:N}^*, u_{1:N-1}^*$ . Now let's look at optimizing over a subsection of this trajectory. That means that instead of solving for  $x_{1:N}, u_{1:N-1}$ , we are now solving for  $x_{L:N}, u_{L:N-1}$  for some new timestep  $1 < L < N$ . What we are going to do is take the initial condition from  $x_L^*$  from our original optimization problem, and setup a new optimization problem that optimizes over  $x_{L:N}, u_{L:N-1}$ :

$$\begin{aligned} \min_{x_{L:N}, u_{L:N-1}} \quad & \sum_{i=L}^{N-1} \left[ \frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \\ \text{st} \quad & x_L = x_L^* \\ & x_{i+1} = A x_i + B u_i \quad \text{for } i = L, L+1, \dots, N-1 \end{aligned}$$

```
In [35]: @testset "Bellman's Principle of Optimality" begin

    # problem setup
    dt = 0.1
    tf = 5.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    A,B = double_integrator_AB(dt)
    nx,nu = size(B)
    x0 = [5,7,2,-1.4] # initial condition
    Q = diagm(ones(nx))
    R = diagm(ones(nu))
    Qf = 5*Q

    # solve for X_{1:N}, U_{1:N-1} with convex optimization
    Xcvx1,Ucvx1 = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)


    # now let's solve a subsection of this trajectory
    L = 18
    N_2 = N - L + 1

    # here is our updated initial condition from the first problem
    x0_2 = Xcvx1[L]
    Xcvx2,Ucvx2 = convex_trajopt(A,B,Q,R,Qf,N_2,x0_2; verbose = false)

    # test if these trajectories match for the times they share
    U_error = Ucvx1[L:end] .- Ucvx2
    X_error = Xcvx1[L:end] .- Xcvx2
    @test 1e-14 < maximum(norm.(U_error)) < 1e-3
    @test 1e-14 < maximum(norm.(X_error)) < 1e-3

    # -----plotting -----
    X1m = mat_from_vec(Xcvx1)
    X2m = mat_from_vec(Xcvx2)
    plot(X2m[1,:],X2m[2,:], label = "optimal subtrajectory", lw = 5, ls = :dot)
    display(plot!(X1m[1,:],X1m[2,:],
                  title = "Trajectory in State Space",
                  ylabel = "x₂", xlabel = "x₁", label = "full trajectory"))
    # -----plotting -----

    @test isapprox(Xcvx1[end], [-0.02285990, -0.07140241, -0.21259, -0.1540299], rtol = 1e-3)
    @test 1e-14 < norm(Xcvx1[end] - Xcvx2[end],Inf) < 1e-3
end
```



Test Summary:	Pass	Total
Bellman's Principle of Optimality	4	4

Out[35]: Test.DefaultTestSet("Bellman's Principle of Optimality", Any[], 4, false, false)

Part C: Finite-Horizon LQR via Ricatti (10 pts)

Now we are going to solve the original finite-horizon LQR problem:

$$\begin{aligned} \min_{x_{1:N}, u_{1:N-1}} \quad & \sum_{i=1}^{N-1} \left[ \frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \\ \text{st} \quad & x_1 = x_{1C} \\ & x_{i+1} = A x_i + B u_i \quad \text{for } i = 1, 2, \dots, N-1 \end{aligned}$$

with a Ricatti recursion instead of convex optimization. We describe our optimal cost-to-go function (aka the Value function) as the following:

$$V_k(x) = \frac{1}{2} x^T P_k x$$

```
In [73]: """
use the Ricatti recursion to calculate the cost to go quadratic matrix P and
optimal control gain K at every time step. Return these as a vector of matrices,
where P_k = P[k], and K_k = K[k]
"""
function fhlqr(A::Matrix, # A matrix
              B::Matrix, # B matrix
              Q::Matrix, # cost weight
              R::Matrix, # cost weight
              Qf::Matrix, # term cost weight
              N::Int64, # horizon size
              )::Tuple{Vector{Matrix{Float64}}, Vector{Matrix{Float64}}} # return two matrices

# check sizes of everything
nx,nu = size(B)
@assert size(A) == (nx, nx)
@assert size(Q) == (nx, nx)
@assert size(R) == (nu, nu)
@assert size(Qf) == (nx, nx)

# instantiate S and K
P = [zeros(nx,nx) for i = 1:N]
K = [zeros(nu,nx) for i = 1:N-1]

# initialize S[N] with Qf
P[N] = deepcopy(Qf)

# Ricatti
for k = N-1:-1:1
    # TODO
    K[k] = (R + B' * P[k+1] * B)\ (B' * P[k+1] * A)
    P[k] = Q + A' * P[k+1] * (A - B * K[k])
end

return P, K
end
```

Out[73]: fhlqr

```
In [75]: @testset "Convex trajopt vs LQR" begin

# problem stuff
dt = 0.1
tf = 5.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
x0 = [5,7,2,-1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 5*Q


# solve for X_{1:N}, U_{1:N-1} with convex optimization
Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)
P, K = fhlqr(A,B,Q,R,Qf,N)
# now let's simulate using Ucvx
Xsim_cvx = [zeros(nx) for i = 1:N]
Xsim_cvx[1] = 1*x0
Xsim_lqr = [zeros(nx) for i = 1:N]
Xsim_lqr[1] = 1*x0
for i = 1:N-1
    # simulate cvx control
    Xsim_cvx[i+1] = A*Xsim_cvx[i] + B*Ucvx[i]

    # TODO: use your FHLQR control gains K to calculate u_lqr
    # simulate lqr control
    u_lqr = -K[i] * Xsim_lqr[i]
    Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr
end

@test isapprox(Xsim_lqr[end], [-0.02286201, -0.0714058, -0.21259, -0.154030], rtol = 1e-3)
@test 1e-13 < norm(Xsim_lqr[end] - Xsim_cvx[end]) < 1e-3
@test 1e-13 < maximum(norm.(Xsim_lqr - Xsim_cvx)) < 1e-3

# -----plotting-----
X1m = mat_from_vec(Xsim_cvx)
X2m = mat_from_vec(Xsim_lqr)
# plot trajectory in x1 x2 space
plot(X1m[1,:],X1m[2,:], label = "cvx trajectory", lw = 4, ls = :dot)
display(plot!(X2m[1,:],X2m[2,:],
              title = "Trajectory in State Space",
              ylabel = "x2", xlabel = "x1", lw = 2, label = "lqr trajectory"))
# -----plotting-----

end
```



Test Summary:		Pass	Total
Convex trajopt vs LQR		3	3

Out[75]: Test.DefaultTestSet("Convex trajopt vs LQR", Any[], 3, false, false)

To emphasize that these two methods for solving the optimization problem result in the same solutions, we are now going to sample initial conditions and run both solutions. You will have to fill in your LQR policy again.



```
In [76]: import Random
Random.seed!(1)
@testset "Convex trajopt vs LQR" begin

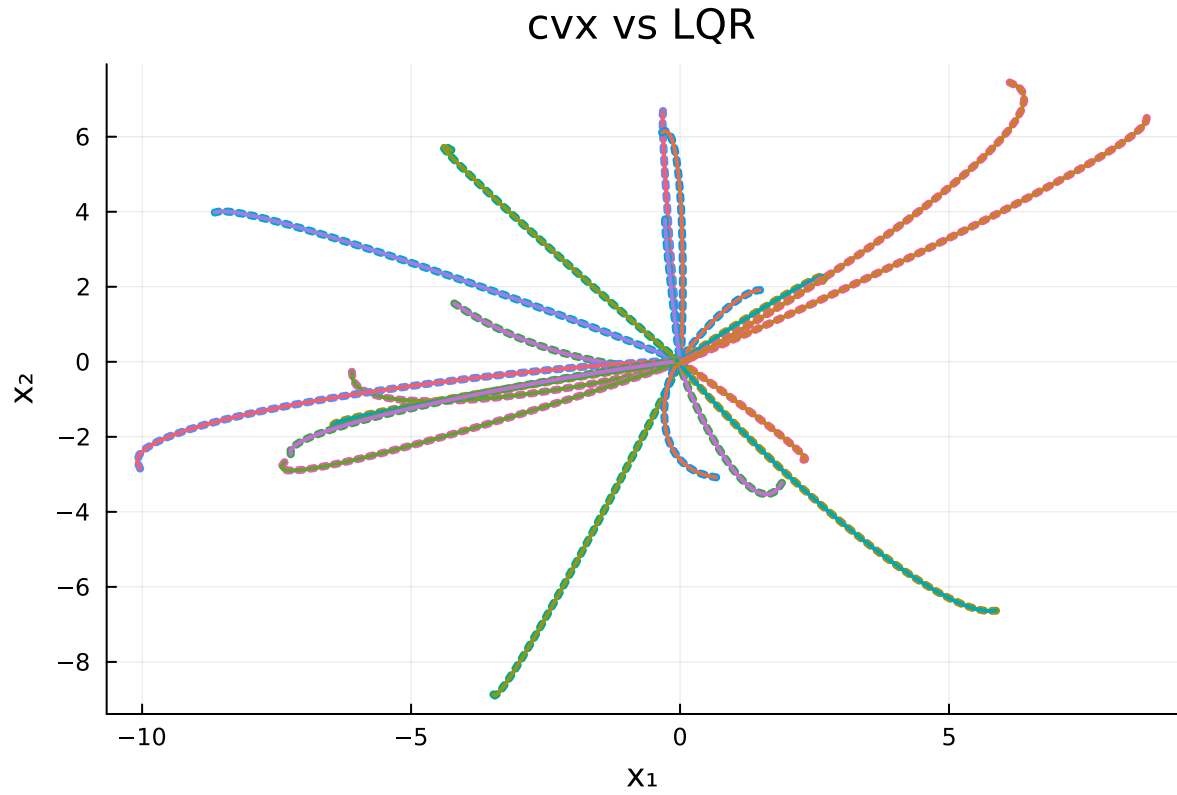
    # problem stuff
    dt = 0.1
    tf = 5.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    A,B = double_integrator_AB(dt)
    nx,nu = size(B)
    Q = diagm(ones(nx))
    R = diagm(ones(nu))
    Qf = 5*Q

    plot()
    for ic_iter = 1:20
        x0 = [5*randn(2); 1*randn(2)]
        # solve for X_{1:N}, U_{1:N-1} with convex optimization
        Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)
        P, K = fhlqr(A,B,Q,R,Qf,N)
        Xsim_cvx = [zeros(nx) for i = 1:N]
        Xsim_cvx[1] = 1*x0
        Xsim_lqr = [zeros(nx) for i = 1:N]
        Xsim_lqr[1] = 1*x0
        for i = 1:N-1
            # simulate cvx control
            Xsim_cvx[i+1] = A*Xsim_cvx[i] + B*Ucvx[i]

            # TODO: use your FHLQR control gains K to calculate u_lqr
            # simulate lqr control
            u_lqr = -K[i] * Xsim_lqr[i]
            Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr
        end

        @test 1e-13 < norm(Xsim_lqr[end] - Xsim_cvx[end]) < 1e-3
        @test 1e-13 < maximum(norm.(Xsim_lqr - Xsim_cvx)) < 1e-3

        # -----plotting-----
        X1m = mat_from_vec(Xsim_cvx)
        X2m = mat_from_vec(Xsim_lqr)
        plot!(X2m[1,:],X2m[2,:], label = "", lw = 4, ls = :dot)
        plot!(X1m[1,:],X1m[2,:], label = "", lw = 2)
    end
    display(plot!(title = "cvx vs LQR", ylabel = "x₂", xlabel = "x₁"))
end
```



Test Summary:	Pass	Total
Convex trajopt vs LQR	40	40

```
Out[76]: Test.DefaultTestSet("Convex trajopt vs LQR", Any[], 40, false, false)
```

Part D: Why LQR is so great (10 pts)

Now we are going to emphasize two reasons why the feedback policy from LQR is so useful:

- 1. It is robust to noise and model uncertainty (the Convex approach would require re-solving of the problem every time the new state differs from the expected state (this is MPC, more on this in Q3)
- 2. We can drive to any achievable goal state with  $u = -K(x - x_{goal})$

First we are going to look at a simulation with the following white noise:

$$x_{k+1} = Ax_k + Bu_k + \text{noise}$$

Where noise  $\sim \mathcal{N}(0, \Sigma)$ .

In [77]: @testset "Why LQR is great reason 1" begin

```
# problem stuff
dt = 0.1
tf = 7.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
x0 = [5,7,2,-1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 10*Q


# solve for X_{1:N}, U_{1:N-1} with convex optimization
Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)
P, K = fhqr(A,B,Q,R,Qf,N)
# now let's simulate using Ucvx
Xsim_cvx = [zeros(nx) for i = 1:N]
Xsim_cvx[1] = 1*x0
Xsim_lqr = [zeros(nx) for i = 1:N]
Xsim_lqr[1] = 1*x0
for i = 1:N-1
    # sampled noise to be added after each step
    noise = [.005*randn(2);.1*randn(2)]

    # simulate cvx control
    Xsim_cvx[i+1] = A*Xsim_cvx[i] + B*Ucvx[i] + noise

    # TODO: use your FHLQR control gains K to calculate u_lqr
    # simulate lqr control
    u_lqr = -K[i] * Xsim_lqr[i]
    Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr + noise
end

# -----plotting-----
X1m = mat_from_vec(Xsim_cvx)
X2m = mat_from_vec(Xsim_lqr)
# plot trajectory in x1 x2 space
plot(X2m[1,:],X2m[2,:], label = "CVX Trajectory (no replanning)", lw = 4, ls = :dot)
display(plot!(X1m[1,:],X1m[2,:],
    title = "Trajectory in State Space (Noisy Dynamics)",
    ylabel = "x2", xlabel = "x1", lw = 2, label = "LQR Trajectory"))
ecvx = [norm(x[1:2]) for x in Xsim_cvx]
elqr = [norm(x[1:2]) for x in Xsim_lqr]
plot(t_vec, elqr, label = "LQR Trajectory",ylabel = "|x - xgoal|",
    xlabel = "time (s)", title = "Error for CVX vs LQR (Noisy Dynamics)")
display(plot!(t_vec, ecvx, label = "CVX Trajectory (no replanning)"))
# -----plotting-----

end
```



Test Summary: |  
Why LQR is great reason 1 | No tests

Out[77]: Test.DefaultTestSet("Why LQR is great reason 1", Any[], 0, false, false)

In [82]: @testset "Why LQR is great reason 2" begin

```
# problem stuff
dt = 0.1
tf = 20.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
x0 = [5,7,2,-1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 10*Q

P, K = fhqr(A,B,Q,R,Qf,N)

# TODO: specify a goal state with 0 velocity within a 5m radius of 0
xgoal = [-3.5,-3.5,0,0]
@test norm(xgoal[1:2])< 5
@test norm(xgoal[3:4])<1e-13 # ensure 0 velocity


Xsim_lqr = [zeros(nx) for i = 1:N]
Xsim_lqr[1] = 1*x0

for i = 1:N-1
    # TODO: use your FHLQR control gains K to calculate u_lqr
    # simulate lqr control
    u_lqr = -K[i] * (Xsim_lqr[i] - xgoal)
    Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr
end

@test norm(Xsim_lqr[end][1:2] - xgoal[1:2]) < .1

# -----plotting-----
Xm = mat_from_vec(Xsim_lqr)
plot(xgoal[1:1],xgoal[2:2],seriestype = :scatter, label = "goal state")
display(plot!(Xm[1,:],Xm[2,:],
               title = "Trajectory in State Space",
               ylabel = "x2", xlabel = "x1", lw = 2, label = "LQR Trajectory"))
```

end



Test Summary:	Pass	Total
Why LQR is great reason 2	3	3

Out[82]: Test.DefaultTestSet("Why LQR is great reason 2", Any[], 3, false, false)

### Part E: Infinite-horizon LQR (10 pts)

Up until this point, we have looked at finite-horizon LQR which only considers a finite number of timesteps in our trajectory. When this problem is solved with a Ricatti recursion, there is a new feedback gain matrix  $K_k$  for each timestep. As the length of the trajectory increases, the first feedback gain matrix  $K_1$  will begin to converge on what we call the "infinite-horizon LQR gain". This is the value that  $K_1$  converges to as  $N \rightarrow \infty$ .

Below, we will plot the values of  $P$  and  $K$  throughout the horizon and observe this convergence.



```
In [80]: # half vectorization of a matrix
function vech(A)
    return A[tril(trues(size(A)))]
end
@testset "P and K time analysis" begin

    # problem stuff
    dt = 0.1
    tf = 10.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    A,B = double_integrator_AB(dt)
    nx,nu = size(B)

    # cost terms
    Q = diagm(ones(nx))
    R = .5*diagm(ones(nu))
    Qf = randn(nx,nx); Qf = Qf'*Qf + I;

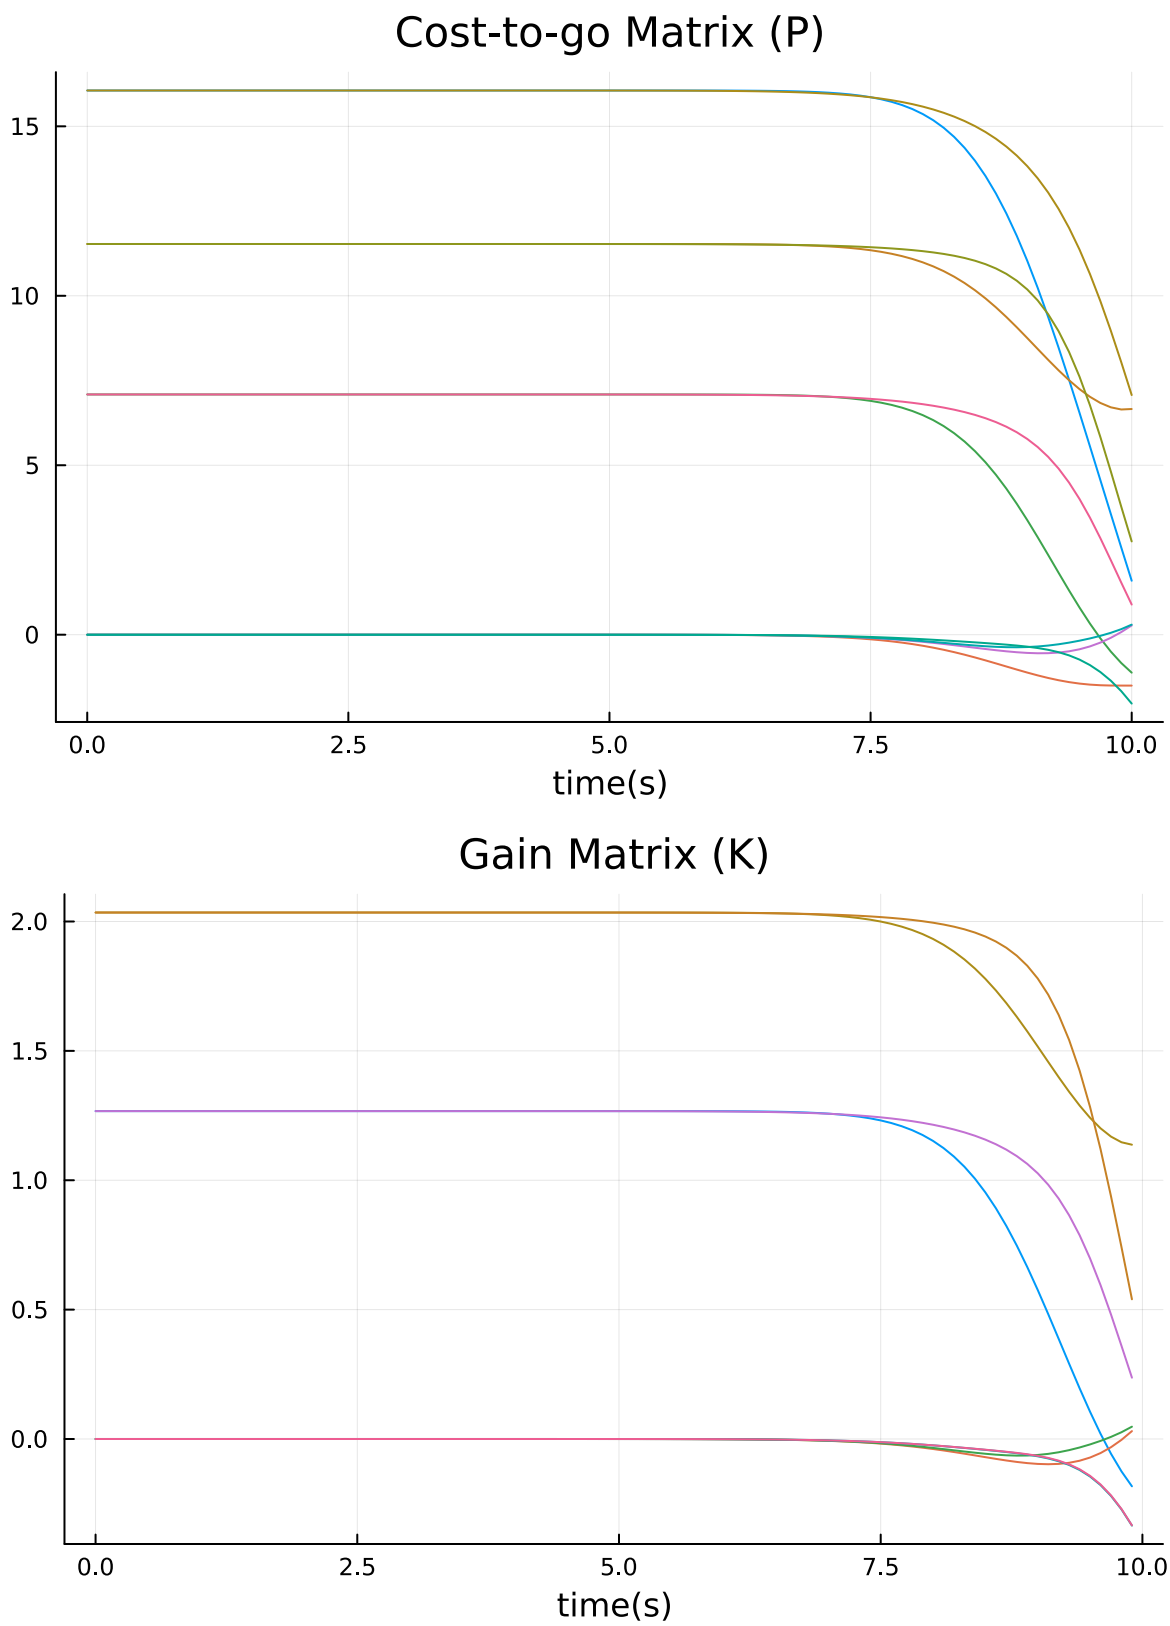
    P, K = fhlqr(A,B,Q,R,Qf,N)

    Pm = hcat(vech.(P)...)
    Km = hcat(vec.(K)...)

    # make sure these things converged
    @test 1e-13 < norm(P[1] - P[2]) < 1e-3
    @test 1e-13 < norm(K[1] - K[2]) < 1e-3

    display(plot(t_vec, Pm', label = "",title = "Cost-to-go Matrix (P)", xlabel = "time(s)"))
    display(plot(t_vec[1:end-1], Km', label = "",title = "Gain Matrix (K)", xlabel = "time(s)"))

end
```



Test Summary:		Pass	Total
P and K time analysis		2	2

Out[80]: Test.DefaultTestSet("P and K time analysis", Any[], 2, false, false)

Complete this infinite horizon LQR function where you do a Ricatti recursion until the cost to go matrix P converges:

$$\|P_k - P_{k+1}\| \leq \text{tol}$$

And return the steady state *P* and *K*.

```
In [88]: """
P,K = ihlqr(A,B,Q,R)

TODO: complete this infinite horizon LQR function where
you do the ricatti recursion until the cost to go matrix
P converges to a steady value |P_k - P_{k+1}| ≤ tol
"""

function ihlqr(A::Matrix,      # vector of A matrices
              B::Matrix,      # vector of B matrices
              Q::Matrix,      # cost matrix Q
              R::Matrix;      # cost matrix R
              max_iter = 1000, # max iterations for Ricatti
              tol = 1e-5      # convergence tolerance
              )::Tuple{Matrix, Matrix} # return two matrices

    # get size of x and u from B
    nx, nu = size(B)

    # initialize S with Q
    P = deepcopy(Q)

    P_prev = 1*P
    for ricatti_iter = 1:max_iter

        K = (R + B' * P_prev * B)\ (B' * P_prev * A)
        P = Q + A' * P_prev * (A - B * K)
        if norm(P - P_prev) <= tol
            @show ricatti_iter
            return P,K
        end
        P_prev = 1*P
    end
    error("ihlqr did not converge")
end
@testset "ihlqr test" begin
    # problem stuff
    dt = 0.1
    A,B = double_integrator_AB(dt)
    nx,nu = size(B)

    # we're just going to modify the system a little bit
    # so the following graphs are still interesting

    Q = diagm(ones(nx))
    R = .5*diagm(ones(nu))
    P, K = ihlqr(A,B,Q,R)

    # check this P is in fact a solution to the Ricatti equation
    @test typeof(P) == Matrix{Float64}
    @test typeof(K) == Matrix{Float64}
    @test 1e-13 < norm(Q + K'*R*K + (A - B*K)'*P*(A - B*K) - P) < 1e-3

end
```

```
ricatti_iter = 68
Test Summary: | Pass Total
ihlqr test   |    3    3
```

Out[88]: Test.DefaultTestSet("ihlqr test", Any[], 3, false, false)

```
In [ ]:
```

```
In [1]: import Pkg
Pkg.activate(@_DIR_)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
import MeshCat as mc
using JLD2
using Test
using Random
include(joinpath(@_DIR_, "utils/cartpole_animation.jl"))

    Activating environment at `~/ocrl_ws/16745-ocrl/HW2_S23/Project.toml`
Precompiling project...
  ✓ Adapt
  ✓ AssetRegistry
  ✓ GPUArraysCore
  ✓ BenchmarkTools
  ✓ GenericLinearAlgebra
  ✓ StructArrays
  ✓ Mux
  ✓ WebIO
  ✓ JSEExpr
  ✓ Blink
  ✓ GeometryBasics
  ✓ JLD2
  ✓ MeshCat
  ✓ MathOptInterface
  ✓ Plots
  ✓ ECOS
  ✓ Convex
  ✓ Hypatia
  18 dependencies successfully precompiled in 78 seconds (173 already precompiled)

Out[1]: animate_cartpole (generic function with 1 method)
```

Q2: LQR for nonlinear systems (25 pts)

Linearization warmup

Before we apply LQR to nonlinear systems, we are going to treat our linear system as if it's nonlinear. Specifically, we are going to "approximate" our linear system with a first-order Taylor series, and define a new set of  $(\Delta x, \Delta u)$  coordinates. Since our dynamics are linear, this approximation is exact, allowing us to check that we set up the problem correctly.

First, assume our discrete time dynamics are the following:

$$x_{k+1} = f(x_k, u_k)$$

And we are going to linearize about a reference trajectory  $\bar{x}_{1:N}, \bar{u}_{1:N-1}$ . From here, we can define our delta's accordingly:

$$\begin{aligned} x_k &= \bar{x}_k + \Delta x_k \\ u_k &= \bar{u}_k + \Delta u_k \end{aligned}$$

Next, we are going to approximate our discrete time dynamics function with the following first order Taylor series:

$$x_{k+1} \approx f(\bar{x}_k, \bar{u}_k) + \left[ \frac{\partial f}{\partial x} \right]_{\bar{x}_k, \bar{u}_k} (x_k - \bar{x}_k) + \left[ \frac{\partial f}{\partial u} \right]_{\bar{x}_k, \bar{u}_k} (u_k - \bar{u}_k)$$

Which we can substitute in our delta notation to get the following:


$$\bar{x}_{k+1} + \Delta x_{k+1} \approx f(\bar{x}_k, \bar{u}_k) + \left[ \frac{\partial f}{\partial x} \right]_{\bar{x}_k, \bar{u}_k} \Delta x_k + \left[ \frac{\partial f}{\partial u} \right]_{\bar{x}_k, \bar{u}_k} \Delta u_k$$

If the trajectory  $\bar{x}, \bar{u}$  is dynamically feasible (meaning  $\bar{x}_{k+1} = f(\bar{x}_k, \bar{u}_k)$ ), then we can cancel these equivalent terms on each side of the above equation, resulting in the following:

$$\Delta x_{k+1} \approx \left[ \frac{\partial f}{\partial x} \right]_{\bar{x}_k, \bar{u}_k} \Delta x_k + \left[ \frac{\partial f}{\partial u} \right]_{\bar{x}_k, \bar{u}_k} \Delta u_k$$

Cartpole

We are now going to look at two different applications of LQR to the nonlinear cartpole system. Given the following description of the cartpole:



(if this image doesn't show up, check out `cartpole.png`)

with a cart position  $p$  and pole angle  $\theta$ . We are first going to linearize the nonlinear discrete dynamics of this system about the point where  $p = 0$ , and  $\theta = 0$  (no velocities), and use an infinite horizon LQR controller about this linearized state to stabilize the cartpole about this goal state. The dynamics of the cartpole are parametrized by the mass of the cart, the mass of the pole, and the length of the pole. To simulate a "sim to real gap", we are going to design our controllers around an estimated set of problem parameters `params_est`, and simulate our system with a different set of problem parameters `params_real`.

```
In [3]: function ihlqr(A::Matrix,      # vector of A matrices
                    B::Matrix,      # vector of B matrices
                    Q::Matrix,      # cost matrix Q
                    R::Matrix;      # cost matrix R
                    max_iter = 1000, # max iterations for Ricatti
                    tol = 1e-5,     # convergence tolerance
                    )::Tuple{Matrix, Matrix} # return two matrices

    # get size of x and u from B
    nx, nu = size(B)

    # initialize S with Q
    P = deepcopy(Q)

    P_prev = 1*P
    for ricatti_iter = 1:max_iter

        K = (R + B' * P_prev * B) \ (B' * P_prev * A)
        P = Q + A' * P_prev * (A - B * K)
        if norm(P - P_prev) <= tol
            @show ricatti_iter
            return P,K
        end
        P_prev = 1*P
    end
    error("ihlqr did not converge")
end

Out[3]: ihlqr (generic function with 1 method)
```

```
In [4]: """
continuous time dynamics for a cartpole, the state is
x = [p, θ, ṗ, θ̇]
where p is the horizontal position, and θ is the angle
where θ = 0 has the pole hanging down, and θ = 180 is up.

The cartpole is parametrized by a cart mass `mc`, pole
mass `mp`, and pole length `l`. These parameters are loaded
into a `params::NamedTuple`. We are going to design the
controller for a estimated `params_est`, and simulate with
`params_real`.
"""
function dynamics(params::NamedTuple, x::Vector, u)
    # cartpole ODE, parametrized by params.

    # cartpole physical parameters
    mc, mp, l = params.mc, params.mp, params.l
    g = 9.81

    q = x[1:2]
    qd = x[3:4]

    s = sin(q[2])
    c = cos(q[2])

    H = [mc+mp mp*l*c; mp*l*c mp*l^2]
    C = [0 -mp*qd[2]*l*s; 0 0]
    G = [0, mp*g*l*s]
    B = [1, 0]

    qdd = -H\C*qd + G - B*u[1]
    return [qd;qdd]

end

function rk4(params::NamedTuple, x::Vector,u,dt::Float64)
    # vanilla RK4
    k1 = dt*dynamics(params, x, u)
    k2 = dt*dynamics(params, x + k1/2, u)
    k3 = dt*dynamics(params, x + k2/2, u)
    k4 = dt*dynamics(params, x + k3, u)
    x + (1/6)*(k1 + 2*k2 + 2*k3 + k4)
end

Out[4]: rk4 (generic function with 1 method)
```

Part A: Infinite Horizon LQR about an equilibrium (10 pts)

Here we are going to solve for the infinite horizon LQR gain, and use it to stabilize the cartpole about the unstable equilibrium.

```
In [37]: @testset "LQR about eq" begin

    # states and control sizes
    nx = 4
    nu = 1

    # desired x and g (linearize about these)
    xgoal = [0, pi, 0, 0]
    ugoal = [0]

    # initial condition (slightly off of our linearization point)
    x0 = [0, pi, 0, 0] + [1.5, deg2rad(-20), .3, 0]

    # simulation size
    dt = 0.1
    tf = 5.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    X = [zeros(nx) for i = 1:N]
    X[1] = x0

    # estimated parameters (design our controller with these)
    params_est = (mc = 1.0, mp = 0.2, l = 0.5)

    # real paremeters (simulate our system with these)
    params_real = (mc = 1.2, mp = 0.16, l = 0.55)

    # TODO: solve for the infinite horizon LQR gain Kinf

    # cost terms
    Q = diagm([1,1,.05,.1])
    Qf = 1*Q
    R = 0.1*diagm(ones(nu))

    A = FD.jacobian(dx->rk4(params_est, dx, ugoal, dt), xgoal)
    B = FD.jacobian(du->rk4(params_est, xgoal, du, dt), ugoal)

    P, Kinf = ihlqr(A,B,Q,R)


    # TODO: simulate this controlled system with rk4(params_real, ...)
    for i = 1:N-1
        X[i+1] = rk4(params_real, X[i], ugoal-Kinf * (X[i]-xgoal), dt)
    end

    # -----tests and plots/animations-----
    @test X[1] == x0
    @test norm(X[end])>0
    @test norm(X[end] - xgoal) < 0.1


    Xm = hcat(X...)
    display(plot(t_vec,Xm',title = "cartpole",
        xlabel = "time(s)", ylabel = "x",
        label = ["p" "θ" "ṗ" "θ̇"]))

    # animation stuff
    display(animate_cartpole(X, dt))
    # -----tests and plots/animations-----
end
```

ricatti\_iter = 74



[ Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:  
http://127.0.0.1:8718



Open Controls

Test Summary:		Pass	Total
LQR about eq		3	3

Out[37]: Test.DefaultTestSet("LQR about eq", Any[], 3, false, false)

Part B: TVLQR for trajectory tracking (15 pts)

Here we are given a swingup trajectory that works for `params_est` , but will fail to work with `params_real` . To account for this sim to real gap, we are going to track this trajectory with a TVLQR controller.





```

# optimized trajectory we are going to try and track
DATA = load(joinpath(@__DIR__, "swingup.jld2"))
Xbar = DATA["X"]
Ubar = DATA["U"]

# states and controls
nx = 4
nu = 1

# problem size
dt = 0.05
tf = 4.0
t_vec = 0:dt:tf
N = length(t_vec)

# states (initial condition of zeros)
X = [zeros(nx) for i = 1:N]
X[1] = [0, 0, 0, 0.0]

# make sure we have the same initial condition
@assert norm(X[1] - Xbar[1]) < 1e-12

# real and estimated params
params_est = (mc = 1.0, mp = 0.2, l = 0.5)
params_real = (mc = 1.2, mp = 0.16, l = 0.55)

# TODO: design a time-varying LQR controller to track this trajectory
# use params_est for your control design, and params_real for the simulation

# cost terms
Q = diagm([1,1,.05,.1])
Qf = 10*Q
R = 0.05*diagm(ones(nu))
A = [zeros(nx,nx) for i = 1:N-1]
B = [zeros(nx,nu) for i = 1:N-1]
K = [zeros(nu,nx) for i = 1:N-1]
P = [zeros(nx,nx) for i = 1:N]

@show size(Xbar)
@show size(Ubar)
@show size(A)
@show size(B)
@show size(K)
@show N

for i = 1:N-1
    # todo ask TAs- which one is okay???
    A[i] = FD.jacobian(dx->rk4(params_est, dx, Ubar[i], dt), Xbar[i+1])
    B[i] = FD.jacobian(du->rk4(params_est, Xbar[i+1], du, dt), Ubar[i])
    A[i] = FD.jacobian(dx->rk4(params_est, dx, Ubar[i], dt), Xbar[i])
    B[i] = FD.jacobian(du->rk4(params_est, Xbar[i], du, dt), Ubar[i])
end

P[N] = deepcopy(Qf)

# Ricatti
for k = N-1:-1:1
    K[k] = (R + B[k]' * P[k+1] * B[k]) \ (B[k]' * P[k+1] * A[k])
    P[k] = Q + A[k]' * P[k+1] * (A[k] - B[k] * K[k])
end

# TODO: simulate this controlled system with rk4(params_real, ...)
for i = 1:N-1
    X[i+1] = rk4(params_real, X[i], Ubar[i] - K[i] * (X[i] - Xbar[i]), dt)
end


# -----tests and plots/animations-----
xn = X[N]
@show xn
@test norm(xn) > 0
@test 1e-6 < norm(xn - Xbar[end]) < .2
@test abs(abs(rad2deg(xn[2])) - 180) < 5 # within 5 degrees


Xm = hcat(X...)
Xbarm = hcat(Xbar...)
plot(t_vec, Xbarm', ls=:dash, label = ["p̄" "θ" "ṗ" "θ̇"], lc = [:red :green :blue :black])
display(plot!(t_vec, Xm', title = "Cartpole TVLQR (-- is reference)",
    xlabel = "time(s)", ylabel = "x",
    label = ["p" "θ" "ṗ" "θ̇"], lc = [:red :green :blue :black]))

# animation stuff
display(animate_cartpole(X, dt))
# -----tests and plots/animations-----

```

Cartpole TVLQR (-- is reference)





Test Summary: | Pass Total  
track swingup | 3 3

Out[39]: Test.DefaultTestSet("track swingup", Any[], 3, false, false)

In [ ]:

```
In [1]: import Pkg
Pkg.activate(@_DIR_)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
import MeshCat as mc
using Test
using Random
import Convex as cvx
import ECOS      # the solver we use in this hw
# import Hypatia # other solvers you can try
# import COSMO   # other solvers you can try
```

Notes:

- 1. Some of the cells below will have multiple outputs (plots and animations), it can be easier to see everything if you do `Cell -> All Output -> Toggle Scrolling` , so that it simply expands the output area to match the size of the outputs.
- 2. Things in space move very slowly (by design), because of this, you may want to speed up the animations when you're viewing them. You can do this in MeshCat by doing `Open Controls -> Animations -> Time Scale` , to modify the time scale. You can also play/pause/scrub from this menu as well.
- 3. You can move around your view in MeshCat by `clicking + dragging` , and you can pan with `right click + dragging` , and zoom with the scroll wheel on your mouse (or trackpad specific alternatives).

```
In [15]: # utilities for converting to and from vector of vectors <-> matrix
function mat_from_vec(X::Vector{Vector{Float64}})::Matrix
    # convert a vector of vectors to a matrix
    Xm = hcat(X...)
    return Xm
end
function vec_from_mat(Xm::Matrix)::Vector{Vector{Float64}}
    # convert a matrix into a vector of vectors
    X = [Xm[:,i] for i = 1:size(Xm,2)]
```

Is LQR the answer for everything?

Unfortunately, no. LQR is great for problems with true quadratic costs and linear dynamics, but this is a very small subset of convex trajectory optimization problems. While a quadratic cost is common in control, there are other available convex cost functions that may better motivate the desired behavior of the system. These costs can be things like an L1 norm on the control inputs ( $\|u\|_1$ ), or an L2 goal error ( $\|x - x_{goal}\|_2$ ). Also, control problems often have constraints like path constraints, control bounds, or terminal constraints, that can't be handled with LQR. With the addition of these constraints, the trajectory optimization problem is still convex and easy to solve, but we can no longer just get an optimal gain  $K$  and apply a feedback policy in these situations.

The solution to this is Model Predictive Control (MPC). In MPC, we are setting up and solving a convex trajectory optimization at every time step, optimizing over some horizon or window into the future, and executing the first control in the solution. To see how this works, we are going to try this for a classic space control problem: the rendezvous.

Q3: Optimal Rendezvous and Docking (50 pts)

In this example, we are going to use convex optimization to control the SpaceX Dragon 1 spacecraft as it docks with the International Space Station (ISS). The dynamics of the Dragon vehicle can be modeled with [Clohessy-Wiltshire equations \(https://en.wikipedia.org/wiki/Clohessy-Wiltshire\\_equations\)](https://en.wikipedia.org/wiki/Clohessy-Wiltshire_equations), which is a linear dynamics model in continuous time. The state and control of this system are the following:

$$x = [r_x, r_y, r_z, v_x, v_y, v_z]^T,$$
$$u = [t_x, t_y, t_z]^T,$$

where  $r$  is a relative position of the Dragon spacecraft with respect to the ISS,  $v$  is the relative velocity, and  $t$  is the thrust on the spacecraft. The continuous time dynamics of the vehicle are the following:

$$\dot{x} = \begin{bmatrix} 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 3n^2 & 0 & 0 & 0 & 2n & 0 \\ 0 & 0 & 0 & -2n & 0 & 0 \\ 0 & 0 & -n^2 & 0 & 0 & 0 \end{bmatrix} x + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} u,$$

where  $n = \sqrt{\mu/a^3}$ , with  $\mu$  being the [standard gravitational parameter \(https://en.wikipedia.org/wiki/Standard\\_gravitational\\_parameter\)](https://en.wikipedia.org/wiki/Standard_gravitational_parameter), and  $a$  being the semi-major axis of the orbit of the ISS.

We are going to use three different techniques for solving this control problem, the first is LQR, the second is convex trajectory optimization, and the third is convex MPC where we will be able to account for unmodeled dynamics in our system (the "sim to real" gap).

Part A: Discretize the dynamics (5 pts)

Use the matrix exponential to convert the linear ODE into a linear discrete time model (hint: the matrix exponential is just `exp()` in Julia when called on a matrix).

```
In [4]: function create_dynamics(dt::Real)::Tuple{Matrix,Matrix}
    mu = 3.986004418e14 # standard gravitational parameter
    a = 6971100.0      # semi-major axis of ISS
    n = sqrt(mu/a^3)   # mean motion

    # continuous time dynamics x-dot = Ax + Bu
    A = [0 0 0 1 0 0;
         0 0 0 0 1 0;
         0 0 0 0 0 1;
         3*n^2 0 0 0 2*n 0;
         0 0 0 -2*n 0 0;
         0 0 -n^2 0 0 0]

    B = Matrix([zeros(3,3);0.1*I(3)])
    nx, nu = size(B)
    # TODO: convert to discrete time X_{k+1} = Ad*x_k + Bd*u_k

    # Ad = zeros(6,6)
    # Bd = zeros(6,3)

    expmat = exp([A B; zeros(nu, nx + nu)] .* dt)

    Ad = expmat[1:nx,1:nx]
    Bd = expmat[1:nx,nx+1:nx+nu]
```

```
In [5]: @testset "discrete dynamics" begin
    A,B = create_dynamics(1.0)

    x = [1,3,-.3,.2,.4,-.5]
    u = [-.1,.5,.3]

    # test these matrices
    @test isapprox(A*x + B*u, [1.195453, 3.424786, -0.78499972, 0.190925, 0.4495759, -0.4699993], atol = 1e-3)
    @test isapprox(det(A), 1, atol = 1e-8)
    @test isapprox(norm(B,Inf), 0.0999999803, atol = 1e-5)

end
```

Part B: LQR (10 pts)

Now we will take a given reference trajectory  $X_{ref}$  and track it with finite-horizon LQR. Remember that finite-horizon LQR is solving this problem:

$$\begin{aligned} \min_{x_1:N,u_1:N-1} \quad & \sum_{i=1}^{N-1} \left[ \frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q_f (x_N - x_{ref,N}) \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \end{aligned}$$

where our policy is  $u_i = -K_i(x_i - x_{ref,i})$ . Use your code from the previous problem with your `fhlqr` function to generate your gain matrices.

One twist we will throw into this is control constraints `u_min` and `u_max`. You should use the function `clamp(u, u_min, u_max)` to clamp the values of your `u` to be within this range.

If implemented correctly, you should see the Dragon spacecraft dock with the ISS successfully, but only after it crashes through the ISS a little bit.

```
In [7]: function fhlqr(A::Matrix, # A matrix
    B::Matrix, # B matrix
    Q::Matrix, # cost weight
    R::Matrix, # cost weight
    Qf::Matrix, # term cost weight
    N::Int64 # horizon size
) :: Tuple{Vector{Matrix{Float64}}, Vector{Matrix{Float64}}} # return two matrices

# check sizes of everything
nx,nu = size(B)
@assert size(A) == (nx, nx)
@assert size(Q) == (nx, nx)
@assert size(R) == (nu, nu)
@assert size(Qf) == (nx, nx)

# instantiate S and K
P = [zeros(nx,nx) for i = 1:N]
K = [zeros(nu,nx) for i = 1:N-1]

# initialize S[N] with Qf
P[N] = deepcopy(Qf)

# Ricatti
for k = N-1:-1:1
    # TODO
    K[k] = (R + B' * P[k+1] * B) \ (B' * P[k+1] * A)
    P[k] = Q + A' * P[k+1] * (A - B * K[k])
end
```



```
In [23]: @testset "LQR rendezvous" begin

# create our discrete time model
dt = 1.0
A,B = create_dynamics(dt)

# get our sizes for state and control
nx,nu = size(B)

# initial and goal states
x0 = [-2;-4;2;0;0;.0]
xg = [0,-.68,3.05,0,0,0]

# bounds on U
u_max = 0.4*ones(3)
u_min = -u_max

# problem size and reference trajectory
N = 120
t_vec = 0:dt:((N-1)*dt)
X_ref = desired_trajectory_long(x0,xg,200,dt)[1:N]

# TODO: FHLQR
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 10*Q
# TODO get K's from fhlqr
P, K = fhlqr(A,B,Q,R,Qf,N)
# simulation
X_sim = [zeros(nx) for i = 1:N]
U_sim = [zeros(nu) for i = 1:N-1]
X_sim[1] = x0
for i = 1:(N-1)
    # TODO: put LQR control law here
    # make sure to clamp
    U_sim[i] = -K[i]* (X_sim[i] - X_ref[i])
    U_sim[i] = -K[i] * X_sim[i]

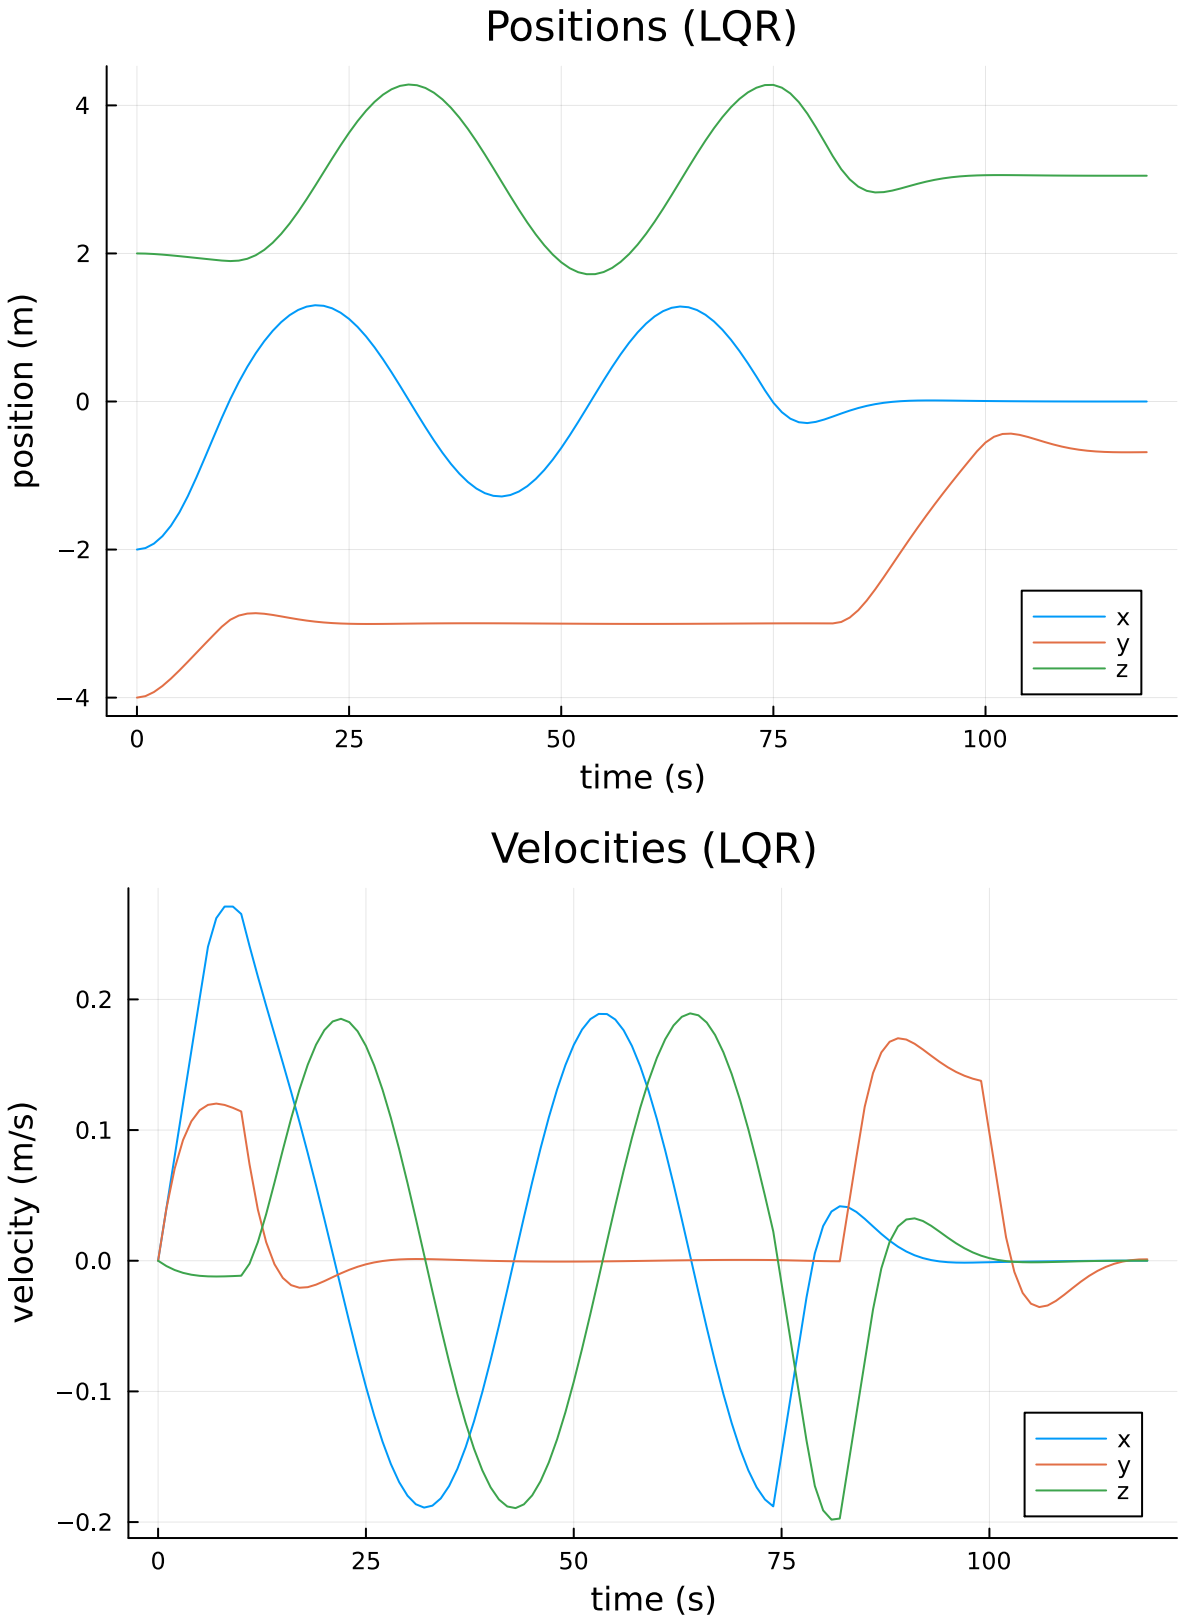
    U_sim[i] = clamp.(U_sim[i], u_min, u_max)

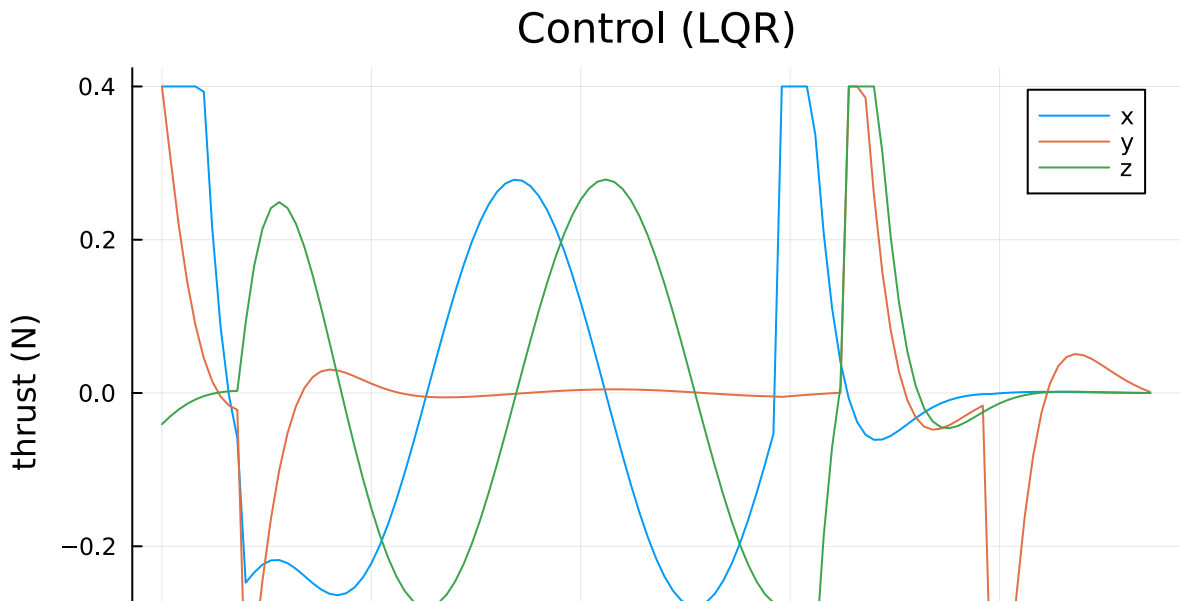
    # simulate 1 step
    X_sim[i+1] = A*X_sim[i] + B*U_sim[i]
end

# -----plotting/animation-----
Xm = mat_from_vec(X_sim)
Um = mat_from_vec(U_sim)
display(plot(t_vec,Xm[1:3,:]',title = "Positions (LQR)",
    xlabel = "time (s)", ylabel = "position (m)",
    label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities (LQR)",
    xlabel = "time (s)", ylabel = "velocity (m/s)",
    label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control (LQR)",
    xlabel = "time (s)", ylabel = "thrust (N)",
    label = ["x" "y" "z"]))


# feel free to toggle `show_reference`
display(animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----

# testing
xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test norm(X_sim[end] - xg) < .01 # goal
@test (xg[2] + .1) < maximum(ys) < 0 # we should have hit the ISS
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_sim)) <= 0.4 # control constraints satisfied
```





**Info:** MeshCat server started. You can open the visualizer by visiting the following URL in your browser:  
<http://127.0.0.1:8727>



Test Summary:	Pass	Total
LQR rendezvous	6	6

Out[23]: Test.DefaultTestSet("LQR rendezvous", Any[], 6, false, false)

### Part C: Convex Trajectory Optimization (15 pts)

Now we are going to assume that we have a perfect model (assume there is no sim to real gap), and that we have a perfect state estimate. With this, we are going to solve our control problem as a convex trajectory optimization problem.

$$\begin{aligned} \min_{x_{1:N}, u_{1:N-1}} \quad & \sum_{i=1}^{N-1} \left[ \frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q_f (x_N - x_{ref,N}) \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = A x_i + B u_i \quad \text{for } i = 1, 2, \dots, N-1 \\ & u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \\ & x_i[2] \leq x_{goal}[2] \quad \text{for } i = 1, 2, \dots, N \\ & x_N = x_{goal} \end{aligned}$$

Where we have an LQR cost, an initial condition constraint ( $x_1 = x_{IC}$ ), linear dynamics constraints ( $x_{i+1} = Ax_i + Bu_i$ ), bound constraints on the control ( $u_{min} \leq u_i \leq u_{max}$ ), an ISS collision constraint ( $x_i[2] \leq x_{goal}[2]$ ), and a terminal constraint ( $x_N = x_{goal}$ ). This problem is convex and we will setup and solve this with `Convex.jl`.

In [30]:

```
"""
Xcvx,Ucvx = convex_trajopt(A,B,X_ref,x0,xg,u_min,u_max,N)

setup and solve the above optimization problem, returning
the solutions X and U, after first converting them to
vectors of vectors with vec_from_mat(X.value)
"""
function convex_trajopt(A::Matrix, # discrete dynamics A
                        B::Matrix, # discrete dynamics B
                        X_ref::Vector{Vector{Float64}}, # reference trajectory
                        x0::Vector, # initial condition
                        xg::Vector, # goal state
                        u_min::Vector, # lower bound on u
                        u_max::Vector, # upper bound on u
                        N::Int64, # length of trajectory
                        )::Tuple{Vector{Vector{Float64}}, Vector{Vector{Float64}}} # return Xcvx,Ucvx

# get our sizes for state and control
nx,nu = size(B)

@assert size(A) == (nx, nx)
@assert length(x0) == nx
@assert length(xg) == nx

# LQR cost
Q = diagm(ones(nx))
R = diagm(ones(nu))

# variables we are solving for
X = cvx.Variable(nx,N)
U = cvx.Variable(nu,N-1)
Xref = mat_from_vec(X_ref)
# TODO: implement cost
obj = 0

for k = 1:(N-1)
    # add stagewise cost
    xk = X[:,k]
    uk = U[:,k]
    xrefk = Xref[:,k]

    obj += (0.5 * cvx.quadform(xk - xrefk,Q))
    obj += (0.5 * cvx.quadform(uk,R))
end
obj += (0.5 * cvx.quadform(X[:,N] - Xref[:,N],Q))

# create problem with objective
prob = cvx.minimize(obj)

# TODO: add constraints with prob.constraints +=
prob.constraints += (X[:,1] == x0)
for k = 1:(N-1)
    # dynamics constraints
    prob.constraints += (X[:,k+1] == A*X[:,k] + B*U[:,k])

    prob.constraints += (U[:,k] <= u_max)
    prob.constraints += (U[:,k] >= u_min)

    prob.constraints += (X[2,k] <= xg[2])
end
prob.constraints += (X[2,N] <= xg[2])
prob.constraints += (X[:,N] == xg)

cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)

X = X.value
U = U.value

Xcvx = vec_from_mat(X)
Ucvx = vec_from_mat(U)

return Xcvx, Ucvx
end

@testset "convex trajopt" begin

    # create our discrete time model
    dt = 1.0
    A,B = create_dynamics(dt)

    # get our sizes for state and control
    nx,nu = size(B)

    # initial and goal states
    x0 = [-2;-4;2;0;0;.0]
    xg = [0,-.68,3.05,0,0,0]

    # bounds on U
    u_max = 0.4*ones(3)
    u_min = -u_max

    # problem size and reference trajectory
    N = 100
    t_vec = 0:dt:((N-1)*dt)
    X_ref = desired_trajectory(x0,xg,N,dt)

    # solve convex trajectory optimization problem
    X_cvx, U_cvx = convex_trajopt(A,B,X_ref, x0,xg,u_min,u_max,N)

    X_sim = [zeros(nx) for i = 1:N]
    X_sim[1] = x0
    for i = 1:N-1
        X_sim[i+1] = A*X_sim[i] + B*U_cvx[i]
    end




    # -----plotting/animation-----
    Xm = mat_from_vec(X_sim)
    Um = mat_from_vec(U_cvx)
    display(plot(t_vec,Xm[1:3,:]',title = "Positions",
                xlabel = "time (s)", ylabel = "position (m)",
                label = ["x" "y" "z"])))
    display(plot(t_vec,Xm[4:6,:]',title = "Velocities",
                xlabel = "time (s)", ylabel = "velocity (m/s)",
                label = ["x" "y" "z"])))
    display(plot(t_vec[1:end-1],Um',title = "Control",
                xlabel = "time (s)", ylabel = "thrust (N)",
                label = ["x" "y" "z"])))

    display(animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
    # -----plotting/animation-----


    @test maximum(norm.( X_sim .- X_cvx, Inf)) < 1e-3
    @test norm(X_sim[end] - xg) < 1e-3 # goal
end
```



```
xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test maximum(ys) <= (xg[2] + 1e-3)
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U cvx.Inf)) <= 0.4 + 1e-3 # control constraints satisfied
```



**Info:** MeshCat server started. You can open the visualizer by visiting the following URL in your browser:  
<http://127.0.0.1:8734>



[Open Controls](#)

**Test Summary:** | **Pass** **Total**  
convex trajopt | **7** **7**

**Out[30]:** Test.DefaultTestSet("convex trajopt", Any[], 7, false, false)

Part D: Convex MPC (20 pts)

In part C, we solved for the optimal rendezvous trajectory using convex optimization, and verified it by simulating it in an open loop fashion (no feedback). This was made possible because we assumed that our linear dynamics were exact, and that we had a perfect estimate of our state. In reality, there are many issues that would prevent this open loop policy from being successful, here are a few:

- imperfect state estimation
- unmodeled dynamics
- misalignments

- actuator uncertainties

Together, these factors result in a "sim to real" gap between our simulated model, and the real model. Because there will always be a sim to real gap, we can't just execute open loop policies and expect them to be successful. What we can do, however, is use Model Predictive Control (MPC) that combines some of the ideas of feedback control with convex trajectory optimization.

A convex MPC controller will set up and solve a convex optimization problem at each time step that incorporates the current state estimate as an initial condition. For a trajectory tracking problem like this rendezvous, we want to track  $x_{ref}$ , but instead of optimizing over the whole trajectory, we will only consider a sliding window of size  $N_{mpc}$  (also called a horizon). If  $N_{mpc} = 20$ , this means our convex MPC controller is reasoning about the next 20 steps in the trajectory. This optimization problem at every timestep will start by taking the relevant reference trajectory at the current window from the current step  $i$ , to the end of the window  $i + N_{mpc} - 1$ . This slice of the reference trajectory that applies to the current MPC window will be called  $\tilde{x}_{ref} = x_{ref}[i, (i + N_{mpc} - 1)]$ .

$$\begin{aligned} \min_{x_{1:N}, u_{1:N-1}} \quad & \sum_{i=1}^{N-1} \left[ \frac{1}{2} (x_i - \tilde{x}_{ref,i})^T Q (x_i - \tilde{x}_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - \tilde{x}_{ref,N})^T Q_f (x_N - \tilde{x}_{ref,N}) \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \\ & u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \\ & x_i[2] \leq x_{goal}[2] \quad \text{for } i = 1, 2, \dots, N \end{aligned}$$

where  $N$  in this case is  $N_{mpc}$ . This allows for the MPC controller to "think" about the future states in a way that the LQR controller cannot. By updating the reference trajectory window ( $\tilde{x}_{ref}$ ) at each step and updating the initial condition ( $x_{IC}$ ), the MPC controller is able to "react" and compensate for the sim to real gap.

You will now implement a function `convex_mpc` where you setup and solve this optimization problem at every timestep, and simply return  $u_1$  from the solution.



In [42]:

```
"""
`u = convex_mpc(A,B,X_ref_window,xic,xg,u_min,u_max,N_mpc)`

setup and solve the above optimization problem, returning the
first control u_1 from the solution (should be a length nu
Vector{Float64}).
"""
function convex_mpc(A::Matrix, # discrete dynamics matrix A
                   B::Matrix, # discrete dynamics matrix B
                   X_ref_window::Vector{Vector{Float64}}, # reference trajectory for this window
                   xic::Vector, # current state x
                   xg::Vector, # goal state
                   u_min::Vector, # lower bound on u
                   u_max::Vector, # upper bound on u
                   N_mpc::Int64, # length of MPC window (horizon)
                   )::Vector{Float64} # return the first control command of the solved policy

    # get our sizes for state and control
    nx,nu = size(B)

    # check sizes
    @assert size(A) == (nx, nx)
    @assert length(xic) == nx
    @assert length(xg) == nx
    @assert length(X_ref_window) == N_mpc

    # LQR cost
    Q = diagm(ones(nx))
    R = diagm(ones(nu))

    # variables we are solving for
    X = cvx.Variable(nx,N_mpc)
    U = cvx.Variable(nu,N_mpc-1)

    # TODO: implement cost function
    obj = 0
    XrefWin = mat_from_vec(X_ref_window)
    for k = 1:(N_mpc-1)
        # add stagewise cost
        xk = X[:,k]

        uk = U[:,k]

        xrefk = XrefWin[:,k]

        obj += (0.5 * cvx.quadform(xk - xrefk,Q))
        obj += (0.5 * cvx.quadform(uk,R))
    end

    obj += (0.5 * cvx.quadform(X[:,N_mpc] - XrefWin[:,N_mpc],Q))

    # create problem with objective
    prob = cvx.minimize(obj)

    # TODO: add constraints with prob.constraints +=
    prob.constraints += (X[:,1] == xic)
    for k = 1:(N_mpc-2)
        # dynamics constraints
        prob.constraints += (X[:,k+1] == A*X[:,k] + B*U[:,k])

        prob.constraints += (U[:,k] <= u_max)
        prob.constraints += (U[:,k] >= u_min)

        prob.constraints += (X[2,k] <= xg[2])
    end

    prob.constraints += (X[2,N_mpc] <= xg[2])
    prob.constraints += (X[:,N_mpc] == xg)

    # solve problem
    cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)

    # get X and U solutions
    X = X.value
    U = U.value

    # return first control U
    return U[:,1]
end

@testset "convex mpc" begin

    # create our discrete time model
    dt = 1.0
    A,B = create_dynamics(dt)

    # get our sizes for state and control
    nx,nu = size(B)

    # initial and goal states
    x0 = [-2;-4;2;0;0;.0]
    xg = [0,-.68,3.05,0,0,0]

    # bounds on U
    u_max = 0.4*ones(3)
    u_min = -u_max

    # problem size and reference trajectory
    N = 100
    t_vec = 0:dt:((N-1)*dt)
    X_ref = [desired_trajectory(x0,xg,N,dt)...,[xg for i = 1:N]...]

    # MPC window size
    N_mpc = 20

    # sim size and setup
    N_sim = N + 20
    t_vec = 0:dt:((N_sim-1)*dt)
    X_sim = [zeros(nx) for i = 1:N_sim]
    X_sim[1] = x0
    U_sim = [zeros(nu) for i = 1:N_sim-1]

    # simulate
    @showprogress "simulating" for i = 1:N_sim-1




        # get state estimate
        xi_estimate = state_estimate(X_sim[i], xg)

        # TODO: given a window of N_mpc timesteps, get current reference trajectory
        X_ref_tilde = X_ref[i:i+N_mpc-1]


        # TODO: call convex mpc controller with state estimate
        u_mpc = convex_mpc(A,B,X_ref_tilde,xi_estimate,xg,u_min,u_max, N_mpc)

        # commanded control goes into thruster model where it gets modified
        U_sim[i] = thruster_model(X_sim[i], xg, u_mpc)
    end
end
```

```
@test maximum(norm.(U sim.Inf)) <= 0.4 + 1e-3 # control constraints
simulating 100% Time: 0:00:02
```



```
http://127.0.0.1:8736
```



Test Summary:	Pass	Total
convex mpc	6	6

Out[42]: Test.DefaultTestSet("convex mpc", Any[], 6, false, false)

In [ ]: